



Frankline
Florence



TREDENCE DATA ENGINEERING Q&A

2025



www.linkedin.com/in/frankline-florence

TREDENCE DATA ENGINEERING INTERVIEW QUESTIONS AND ANSWERS- 2025

Transformations & Logic

1. What's the difference between `groupByKey()` and `reduceByKey()`?
2. Can you walk through a DAG for a job that uses `groupBy` + `join` + `withColumn`?
3. What causes a shuffle? How can you identify and reduce it?
4. Why is `coalesce()` preferred over `repartition()` before writing data?

Delta Lake + Merge Strategy

5. How would you handle schema evolution in a Delta table with streaming input?
6. What happens when two merge operations run on the same Delta table?
7. How do you implement audit columns (`created_dt`, `updated_dt`) in Delta merge logic?
8. What is the role of the `_delta_log` and how is ACID enforced?

Memory Tuning & Debugging

9. Your job is failing with "GC Overhead Limit Exceeded" – what steps will you take?
10. How do you monitor task skew in Spark UI and fix it?
11. How does `persist()` differ from `cache()` – and which is safer for streaming pipelines?
12. What's the impact of setting `spark.sql.shuffle.partitions` too high?

Streaming Scenarios

13. How do you design a watermark strategy for late-arriving Kafka events?
14. What's the difference between `trigger(once)` and `trigger(processingTime)`?
15. How do you recover from checkpoint corruption in structured streaming?

ANSWERS

⚙️ TRANSFORMATIONS AND LOGIC

■ What's the difference between `groupByKey()` and `reduceByKey()`?

- `groupByKey()`:
 - Groups all values with the same key into a **single list** on the same partition.
 - Causes **heavy data shuffling**, which is inefficient for large datasets.
 - Not preferred for aggregations due to its memory overhead.
- `reduceByKey()`:
 - Performs **local aggregation** (on each partition) **before** shuffling.
 - More efficient and **saves network I/O**.
 - Recommended for large-scale aggregations.

➡ Use `reduceByKey()` for aggregations whenever possible.

■ Can you walk through a DAG for a job that uses `groupBy` + `join` + `withColumn`?

1. `groupBy`:

- Triggers a shuffle to group records by key across partitions.
- Outputs an intermediate logical plan node.

2. `join`:

- If not a broadcast join, also triggers a shuffle on both sides.
- Depending on the join type (e.g., inner, left), the DAG will have additional stages for preparing and joining datasets.

3. `withColumn`:

- A narrow transformation; adds or transforms columns.
- Does not trigger a shuffle.

DAG Summary:

Read -> Shuffle (`groupBy`) -> Shuffle (`join`) -> Narrow (`withColumn`) -> Write

What causes a shuffle? How can you identify and reduce it?

Shuffle occurs when:

- Data moves across partitions or executors.
- Common causes:
 - `groupBy`, `join`, `distinct`, `repartition`, `coalesce` (in some cases), `sort`.

How to identify:

- In Spark UI → DAG or Stage view → Look for "Exchange" nodes or long shuffle read/write times.

How to reduce:

- Use `reduceByKey` instead of `groupByKey`.
 - Use broadcast joins for smaller tables.
 - Avoid unnecessary `repartition()`.
 - Use partition pruning and bucketing where applicable.
-

Why is `coalesce()` preferred over `repartition()` before writing data?

`coalesce(n)`:

- Reduces number of partitions without **shuffle**.

- Merges adjacent partitions.
- Efficient for **narrow transformations** like write optimization.

repartition(n):

- **Shuffles data** to create evenly sized partitions.
- Expensive and time-consuming.

Use Case:

- Before writing output (like to S3 or HDFS), use **coalesce()** to reduce small files and optimize I/O.

DELTA LAKE + MERGE STRATEGY

How would you handle schema evolution in a Delta table with streaming input?

- Enable **schema evolution options**:

```
.option("mergeSchema", "true") # For batch
```

```
.option("cloudFiles.schemaEvolutionMode", "rescue") # For Auto Loader
```

- For structured streaming:

Use **Auto Loader (cloudFiles)** or set

```
spark.databricks.delta.schema.autoMerge.enabled = true.
```

Best Practices:

- Handle schema changes explicitly.
- Use **schema enforcement** + evolution cautiously to avoid silent data corruption.

What happens when two merge operations run on the same Delta table?

- Delta Lake uses **optimistic concurrency control**:
 - Both merges will read the same snapshot.
 - First write commits; the second fails with a **conflict**.
 - The second must **retry** using the new snapshot.

Best Practice:

- Serialize merge jobs or use **isolationLevel** with **MERGE** to avoid write conflicts.
-

■ How do you implement audit columns (**created_dt**, **updated_dt**) in Delta merge logic?

Use conditional logic in **MERGE** statement:

In SQL:

```
MERGE INTO target t
USING source s
ON t.id = s.id
WHEN MATCHED THEN UPDATE SET
    t.name = s.name,
    t.updated_dt = current_timestamp()
WHEN NOT MATCHED THEN INSERT *
    VALUES (s.id, s.name, current_timestamp(), current_timestamp())
```

In PySpark:

```
merge_condition = "t.id = s.id"
deltaTable.alias("t").merge(
    sourceDF.alias("s"),
    merge_condition
).whenMatchedUpdate(set={
    "name": "s.name",
    "updated_dt": "current_timestamp()"
}).whenNotMatchedInsert(values={
    "id": "s.id",
    "name": "s.name",
    "created_dt": "current_timestamp()",
    "updated_dt": "current_timestamp()"
}).execute()
```

■ What is the role of the `_delta_log` and how is ACID enforced?

- `_delta_log/`:
 - Stores Delta Lake **transaction logs** as JSON + parquet files.
 - Each file represents a commit.
- **ACID is enforced** through:
 - **Atomic commits**: Files written first, then log is committed.
 - **Isolation**: Snapshots are consistent; readers never see partial writes.
 - **Concurrency control**: Writes use optimistic concurrency with conflict detection.
 - **Durability**: Log files can rebuild table state.

MEMORY TUNING AND DEBUGGING

Your job is failing with “GC Overhead Limit Exceeded” – what steps will you take?

Cause:

- JVM spends too much time in garbage collection (GC) with little memory reclaimed.

Fixes:

1. Increase executor memory:

```
--executor-memory 4g
```

2. Reduce data per executor:

- Use `repartition()`, filter early.

3. Use Kryo serializer:

```
spark.conf.set("spark.serializer",  
"org.apache.spark.serializer.KryoSerializer")
```

4. Persist selectively and avoid caching large datasets unnecessarily.

5. Tune GC settings (if using JVM directly).

6. Use broadcast joins or DataFrame optimizations.

How do you monitor task skew in Spark UI and fix it?

Detect Skew:

- Spark UI → Stages tab → Look for tasks with:
 - Very long execution time.
 - Large input size compared to others.

Fixes:

- **Salting keys:** Add random prefixes to skewed keys.
- **Broadcast small tables** to avoid join shuffle.
- Use **skew hints** in Spark 3.x:

```
SELECT /*+ SKew('key') */ ...
```

- Use **adaptive execution** (enabled by default in Spark 3+).
-

■ How does `persist()` differ from `cache()` – and which is safer for streaming pipelines?

`cache()`:

- Shortcut for `persist(StorageLevel.MEMORY_AND_DISK)`.

`persist(level)`:

- Offers **fine-grained control**: `MEMORY_ONLY`, `MEMORY_AND_DISK`, `DISK_ONLY`, etc.

In Streaming:

- `persist()` with `MEMORY_AND_DISK_SER` or `DISK_ONLY` is safer:
 - Prevents OOM (Out Of Memory) errors in long-running jobs.
 - Avoids GC pressure.
-

■ What's the impact of setting `spark.sql.shuffle.partitions` too high?

Too many partitions:

- More **task overhead** (scheduler burden).
- High **disk I/O** and **small files**.
- **Increased shuffle latency**.

Too few:

- May **underutilize cluster resources**.
- Large partitions = **skewed tasks**, OOM errors.

Best Practice:

- Tune it based on data size:
 - 200–500 for large jobs.
 - Use `df.rdd.getNumPartitions()` to inspect.
 - Use **adaptive execution** to dynamically tune at runtime.

■ How do you design a watermark strategy for late-arriving Kafka events?

Use `withWatermark(event_time, "delay threshold")`:

```
>> df.withWatermark("event_time", "10 minutes")
```

- Defines how long to **wait for late data**.
- Late events **older than watermark** are **dropped** during aggregation.

Design Tips:

- Choose threshold based on:
 - Expected event delays.
 - Tolerance for late/missing data.
 - State retention limits.

■ What's the difference between `trigger(once)` and `trigger(processingTime)`?

`trigger(once)`:

- Runs **once**, processes all available data, then stops.
- Good for **micro-batch jobs that mimic batch processing**.

`trigger(processingTime="5 seconds")`:

- Runs **every interval**, checking for new data.
 - Keeps streaming jobs alive
-

■ How do you recover from a checkpoint corruption in structured streaming?

1. Stop the streaming job.
2. Backup corrupted checkpoint (for inspection).
3. Remove the checkpoint directory.
4. Restart the job with:
 - Clean checkpoint.
 - Optionally use:

```
>> .option("startingOffsets", "earliest") # Kafka
```
5. Consider using **Delta Lake for idempotent sink** to prevent duplicates during reprocessing.