

◆ Beginner Level PySpark Coding Questions with Answers

1 Read CSV file into DataFrame

```
df = spark.read.csv("path/to/file.csv", header=True, inferSchema=True)
```

✓ **Explanation:** Reads CSV with headers and infers data types automatically.

2 Select specific columns from a DataFrame

```
df.select("column1", "column2").show()
```

✓ **Explanation:** Selects and displays specific columns from the DataFrame.

3 Filter rows based on a condition

```
df.filter(df["age"] > 30).show()
```

✓ **Explanation:** Filters rows where the age is greater than 30.

4 Add a new column using withColumn()

```
df = df.withColumn("new_col", df["salary"] * 0.1)
```

✓ **Explanation:** Adds a new column new_col as 10% of the salary.

5 Group by and aggregate

```
df.groupBy("department").agg({"salary": "avg"}).show()
```

✓ **Explanation:** Computes the average salary for each department.

◆ Intermediate PySpark Coding Questions with Answers

6 Join two DataFrames

```
joined_df = df1.join(df2, df1.id == df2.id, "inner")
```

✓ **Explanation:** Performs an inner join on the id column between two DataFrames.

7 Drop duplicates

```
df.dropDuplicates(["name", "age"]).show()
```

✅ **Explanation:** Removes rows that have the same combination of name and age.

8 Handle null values

```
df.na.fill("Unknown").na.drop()
```

✅ **Explanation:** Fills nulls with "Unknown" and then drops any remaining nulls.

9 Window functions (e.g., row_number, rank)

```
from pyspark.sql.window import Window
```

```
from pyspark.sql.functions import row_number
```

```
window_spec = Window.partitionBy("department").orderBy("salary")
```

```
df.withColumn("row_num", row_number().over(window_spec)).show()
```

✅ **Explanation:** Assigns a row number to each row within a department based on salary order.

10 Convert string to timestamp

```
from pyspark.sql.functions import to_timestamp
```

```
df = df.withColumn("timestamp_col", to_timestamp("date_col", "yyyy-MM-dd HH:mm:ss"))
```

✅ **Explanation:** Converts the date_col string to timestamp format.

🔥 Advanced PySpark Coding Questions with Answers

1 1 Explode a nested array column

```
from pyspark.sql.functions import explode
```

```
df = df.withColumn("exploded_col", explode(df["array_col"]))
```

✅ **Explanation:** Flattens an array column into multiple rows.

1 2 Pivot data

```
df.groupBy("name").pivot("year").agg({"sales": "sum"}).show()
```

✓ **Explanation:** Converts unique year values into separate columns with summed sales.

1 3 Create and use UDF (User Defined Function)

```
from pyspark.sql.functions import udf
```

```
from pyspark.sql.types import StringType
```

```
def get_length(s):
```

```
    return len(s)
```

```
length_udf = udf(get_length, StringType())
```

```
df = df.withColumn("length", length_udf(df["name"]))
```

✓ **Explanation:** Defines a UDF to calculate string length and applies it to the name column.

1 4 Broadcast join for performance

```
from pyspark.sql.functions import broadcast
```

```
df = df1.join(broadcast(df2), "id")
```

✓ **Explanation:** Broadcasts the smaller DataFrame to all executors to avoid shuffles during joins.

1 5 Read from and write to Parquet

```
# Write to Parquet
```

```
df.write.parquet("path/to/output.parquet")
```

```
# Read from Parquet
```

```
df_parquet = spark.read.parquet("path/to/output.parquet")
```

✓ **Explanation:** Saves the DataFrame as a Parquet file and reads it back.

Performance & Optimization Questions with Answers

1 6 How to cache or persist a DataFrame

```
df.cache()    # Stores in memory
```

```
df.persist()  # Can store in memory or disk (customizable)
```

✓ **Explanation:** Avoids recomputation by storing the DataFrame in memory or disk.

1 7 How to check the execution plan

```
df.explain(True)
```

✓ **Explanation:** Shows a detailed physical execution plan, helping diagnose issues like shuffles or scans.

1 8 Partitioning a DataFrame before saving

```
df.write.partitionBy("year", "month").parquet("path")
```

✓ **Explanation:** Saves data into folders based on year and month to optimize query performance.

Scenario-Based Questions with Solutions

1 9 Find top N records per group (e.g., top 3 salaries per department)

```
from pyspark.sql.window import Window
```

```
from pyspark.sql.functions import row_number
```

```
windowSpec = Window.partitionBy("department").orderBy(df["salary"].desc())
```

```
df = df.withColumn("rank", row_number().over(windowSpec)).filter("rank <= 3")
```

✓ **Explanation:** Assigns a rank based on salary within each department and filters top 3.

2 0 Detect duplicates and count them

```
df.groupBy("id").count().filter("count > 1").show()
```

✅ **Explanation:** Groups by id, counts occurrences, and filters where count is greater than 1 (duplicates).

Bonus – Common Conceptual Questions:

- What is the difference between transformations and actions in Spark?
 - Explain wide vs. narrow transformations.
 - What causes shuffles, and how do they impact performance?
 - How does Spark handle fault tolerance?
 - Explain Spark's DAG and stages.
-

✅ **Conclusion:**

Mastering these PySpark coding patterns not only prepares you for interviews but also makes you job-ready for real-world data engineering challenges on **Databricks, AWS, Azure, or GCP.** 🙌 🚀