

🔥 1 Executor Memory vs. Driver Memory

👉 What's the difference?

- **Driver Memory:**

The driver is the **master node** that runs the Spark application, handles the **DAG scheduling**, tracks tasks, and manages metadata. Driver memory is used for things like Spark context, job coordination, task tracking, and data structures collected back to the driver (e.g., `.collect()` or `.take()`).

- **Executor Memory:**

Executors are the **worker nodes** that run the tasks and perform computations. Executor memory is used for processing data (transformations/actions), storing shuffle data, and caching/persisted datasets (`.cache()` or `.persist()`).

👉 Impact on Garbage Collection, Job Stability, and Performance:

- If **driver memory is too low**, the application can fail due to **OutOfMemory errors** when collecting too much data to the driver (common with `.collect()` misuse).
 - If **executor memory is insufficient**, tasks may fail due to **OOM during shuffles, joins, aggregations, or caching**.
 - **Large executor memory can cause longer garbage collection (GC) pauses**, affecting stability. It's sometimes better to have **more executors with less memory each** to reduce GC overhead.
 - Poor memory configuration leads to job retries, slowdowns, or even crashes.
-

🔥 2 Databricks Clusters – Job vs. Standard

👉 How do you decide?

- **Job Cluster:**

- Spin-up just for a single job or notebook run and then terminates.
- Used for **production jobs, scheduled tasks, or CI/CD pipelines**.
- **Cost-effective** because it auto-terminates when the job finishes.

- **Standard Cluster:**

- Long-running interactive cluster.
- Used for **ad-hoc analysis, development, and collaborative notebooks**.
- Supports multiple users and multiple jobs at once.

👉 Config for a 200GB Pipeline:

- **Cluster Type:** Job cluster for ETL/Batch jobs.
 - **Worker Type:** Choose **compute-optimized (like Standard_DS3_v2)** or memory-optimized based on transformation types.
 - **Autoscaling:**
 - Enable autoscaling between **min 4 to max 12 workers** (depending on job parallelism needs).
 - **Photon:**
 - Enable **Photon acceleration** if the workload involves SQL-heavy transformations — it provides massive speed improvements for SQL and Delta workloads.
 - **Driver:** Medium-size driver (matching worker specs unless driver-heavy operations are involved).
-

🔥 3 Dynamic Allocation – ON or OFF?

👉 When does it actually hurt performance?

- **Dynamic Allocation ON:**
 - Scales executors up/down based on workload.
 - Hurts when tasks require **persistent parallelism**, like **large shuffles, broadcast joins, or streaming jobs**, where releasing executors frequently leads to recomputing cached shuffle data.
- **Dynamic Allocation OFF:**
 - Better for jobs with **large joins, shuffles, or stages that maintain high load throughout**.
 - Avoids frequent executor loss, reduces shuffle overhead, and improves stability.

👉 Trade-offs:

- **ON:** Saves costs for idle time, but can introduce shuffle data loss and recomputation.
 - **OFF:** More stable for large pipelines but keeps the cluster size fixed, leading to higher costs if not tuned properly.
-

🔥 4 Skewed Joins

👉 What causes them?

- When one or more keys in a join have **significantly more data** than others, leading to an uneven distribution of work across tasks. Some tasks finish quickly, while others take a long time — causing performance bottlenecks.

👉 Solutions:

- **Salting:**
 - Add a random number (salt) to skewed keys to spread them across partitions, then remove the salt after the join.
 - **Broadcast Join:**
 - Broadcast the smaller dataset to all executors, avoiding shuffles entirely.
 - Works when one table fits within **driver/executor memory (~2GB limit typically)**.
 - **Repartitioning:**
 - Use `repartition()` or skew join hints to redistribute the skewed data evenly before join operations.
 - **Skew Join Hints (Databricks/Spark 3.0+):**
 - Use `JOIN /*+ SKEW */` syntax to let Spark handle skewed keys automatically.
-

🔥 5 Spark DAG & UI Bottleneck Analysis

👉 How does DAG visualization help debugging?

- **DAG (Directed Acyclic Graph)** shows how Spark breaks a job into stages and tasks.
- Helps identify:
 - Expensive shuffles (shown as stage boundaries).
 - Unnecessary wide transformations (`groupByKey` vs. `reduceByKey`).
 - Long-running stages due to data skew or lack of parallelism.
 - Whether caching/persisting is working effectively.

👉 Key Spark UI Metrics to Check First:

1. **Stage Duration:** Look for stages taking much longer than others.

2. **Task Time Variance:** Check if some tasks are slow (possible skew).
 3. **Shuffle Read/Write Size:** Large shuffle data can indicate poorly planned joins or aggregations.
 4. **Executor CPU/Memory Utilization:** See if tasks are compute-bound or memory-bound.
 5. **GC Time:** High garbage collection time affects job throughput.
-

Final Thoughts:

If you understand these deeply, you are not just a Spark user — you're an optimizer. 🧠 🔥