

1. How would you implement Slowly Changing Dimensions (SCD Type 2 vs Type 3)? Explain the differences and when to use each.

✓ Answer:

SCD Type 2 (Full History Tracking):

- Maintains historical data by inserting a **new record** each time an attribute changes.
- Requires tracking **effective\_date**, **end\_date**, and a **current\_flag** or **version number**.

📌 When to use:

When you need a **complete audit trail** of changes, such as address changes over time.

sql

CopyEdit

-- Example Structure:

customer_id	name	address	start_date	end_date	is_current
-------------	------	---------	------------	----------	------------

-----	-----	-----	-----	-----	-----
-------	-------	-------	-------	-------	-------

123	John	NY	2023-01-01	2024-01-10	0
-----	------	----	------------	------------	---

123	John	LA	2024-01-11	NULL	1
-----	------	----	------------	------	---

---

SCD Type 3 (Limited History Tracking):

- Stores **previous values** in **additional columns** of the same row.
- Typically supports **only one level of historical change**.

📌 When to use:

When only the **current and previous values** are required for reference, like last and current department.

sql

CopyEdit

-- Example Columns:

employee_id	name	current_department	previous_department
-------------	------	--------------------	---------------------

---

2. Describe the differences between transient tables and temporary tables in data warehousing.

✓ **Answer:**

Feature	Temporary Table	Transient Table
Lifetime	Session-level (auto-drops on session end)	Persistent (until manually dropped)
Storage	Cache / Memory	Disk
Fail-safe & Time Travel	Not supported	Not supported
Use Case	For temporary/adhoc logic during pipelines	Intermediate staging without long-term retention

✚ **Summary:**

- Use **temporary tables** for lightweight, in-session operations.
- Use **transient tables** when you want persistence without incurring fail-safe storage costs (e.g., in Snowflake).

---

**3. Write an SQL query using window functions to find the most recent designation of each employee.**

✓ **Answer:**

sql

CopyEdit

```
SELECT *
```

```
FROM (
```

```
    SELECT *,
```

```
        ROW_NUMBER() OVER (PARTITION BY employee_id ORDER BY update_date DESC) AS  
rn
```

```
    FROM employee_designation
```

```
) t
```

```
WHERE rn = 1;
```

✚ **Explanation:**

- The ROW\_NUMBER() window function ranks rows within each employee\_id group by update\_date.
  - We filter for the latest (rn = 1) to get the most recent designation.
- 

#### 4. Explain how you'd model a “swipe payment” API in a relational schema.

✅ Answer:

To represent a **swipe payment system**:

##### Tables:

- **Users:** user\_id, name, email, created\_at
- **Cards:** card\_id, user\_id, card\_type, status
- **Transactions:** txn\_id, card\_id, user\_id, amount, location, swipe\_time, status, device\_id

##### Relationships:

- One-to-Many from Users → Cards
- One-to-Many from Cards → Transactions

📌 Notes:

- Add indexing on user\_id, swipe\_time for querying performance.
  - Add foreign keys to maintain data integrity.
- 

#### 5. What's the difference between Change Data Capture (CDC) and Change Data Tracking (CDT)?

✅ Answer:

Feature	CDC (Change Data Capture)	CDT (Change Data Tracking)
Granularity	Captures actual <b>INSERT/UPDATE/DELETE</b> data	Tracks only <b>that a change occurred</b>
Before/After Values	Supported (can track what changed)	Usually <b>not supported</b>
Use Case	Replication, streaming pipelines	Audit, triggers for lightweight change logic

Feature	CDC (Change Data Capture)	CDT (Change Data Tracking)
Complexity	Requires change log / version tracking	Simpler implementation (e.g., change flags)

✦ **CDC** is used for **data replication** and streaming (e.g., Kafka, Debezium), while **CDT** is suitable for systems that just need to react to changes without knowing the actual data differences.

## 6. Compare data lake storage vs blob storage. When would you use one over the other?

✓ **Answer:**

Feature	Blob Storage	Data Lake Storage (ADLS Gen2)
Hierarchical Namespace	✗ No	✓ Yes (folders, directories)
Optimized for	Unstructured object storage	Big data analytics & high throughput workloads
Security	Basic (limited to container-level)	Advanced (POSIX permissions, RBAC, ACLs)
Performance	Suitable for general-purpose use	Optimized for analytics engines (like Spark, Hive)
Integration	General storage (images, docs, backups)	Tight integration with Azure analytics tools

✦ **When to use:**

- Use **Blob** for general file storage.
- Use **ADLS Gen2** for analytics workloads, hierarchical file systems, and scenarios requiring fine-grained access control and scalable parallel reads/writes.

---

## 7. Walk me through how you'd design an end-to-end ETL pipeline from an on-premise database to Azure Databricks.

✓ **Answer:**

### ◆ ETL Pipeline Design (Step-by-Step):

#### 1. Extract:

- Use **Self-hosted IR** in **Azure Data Factory (ADF)** to connect to on-premise SQL Server, Oracle, or other DB.
- Use ADF **Copy Activity** to pull data.

## 2. Land (Stage):

- Store extracted data as CSV/Parquet in **ADLS Gen2** (Bronze Layer of Data Lake).

## 3. Transform:

- Trigger an **ADF pipeline** to run a **Databricks notebook** using Databricks activity.
- In the notebook, use **PySpark** to clean, transform, deduplicate, join data.

## 4. Load:

- Write transformed data to curated zones in **Delta format** (Silver/Gold Layer) in ADLS.
- Optionally, load aggregated data into **Synapse or Power BI** datasets for visualization.

✦ Add **monitoring** via ADF alerts or logs.

---

**8. In Python, how do you extract unique values from a dictionary? Provide a code snippet.**

✅ **Answer:**

If you want to extract unique values from a dictionary:

python

CopyEdit

```
my_dict = {  
    'a': 10,  
    'b': 20,  
    'c': 10,  
    'd': 30  
}
```

```
unique_values = set(my_dict.values())  
print(unique_values) # Output: {10, 20, 30}
```

✦ `set()` removes duplicates from dictionary `.values()`.

---

## 9. How would you execute one notebook from another in Databricks?

✅ **Answer:**

You can use `dbutils.notebook.run()`:

python

CopyEdit

# In Notebook A

```
dbutils.notebook.run("/path/to/NotebookB", timeout_seconds=60, arguments={"input":  
"value"})
```

✦ Parameters:

- **Path:** Relative or absolute path to target notebook
- **Timeout:** In seconds
- **Arguments:** Passed as dictionary, accessed in the target notebook using `dbutils.widgets.get("key")`

✦ This is helpful for orchestrating modular pipelines.

---

## 10. Write Python code to sum amount per customer\_id given a DataFrame or list of records.

✅ **Answer:**

python

CopyEdit

```
import pandas as pd
```

```
data = [  
    {'customer_id': 1, 'amount': 100},
```

```
{'customer_id': 2, 'amount': 200},  
{'customer_id': 1, 'amount': 50},  
{'customer_id': 2, 'amount': 150}  
]
```

```
df = pd.DataFrame(data)
```

```
result = df.groupby('customer_id')['amount'].sum().reset_index()  
print(result)
```

✦ **Output:**

nginx

CopyEdit

```
customer_id amount  
0          1  150  
1          2  350
```

Alternatively, in PySpark:

python

CopyEdit

```
from pyspark.sql.functions import sum
```

```
df.groupBy("customer_id").agg(sum("amount").alias("total_amount")).show()
```

---

**11. Given a Boolean matrix, count the number of islands (connected components). How would you solve it?**

✅ **Answer:**

Use **Depth-First Search (DFS)** or **Breadth-First Search (BFS)**.

python

CopyEdit

```

def num_islands(grid):
    if not grid:
        return 0

    def dfs(r, c):
        if r < 0 or r >= len(grid) or c < 0 or c >= len(grid[0]) or grid[r][c] != 1:
            return
        grid[r][c] = -1 # mark as visited
        dfs(r+1, c)
        dfs(r-1, c)
        dfs(r, c+1)
        dfs(r, c-1)

    count = 0
    for r in range(len(grid)):
        for c in range(len(grid[0])):
            if grid[r][c] == 1:
                dfs(r, c)
                count += 1
    return count

# Example usage:
matrix = [
    [1, 1, 0],
    [0, 1, 0],
    [0, 0, 1]
]

print(num_islands(matrix)) # Output: 2

```



📌 An "island" is a group of adjacent 1s (connected horizontally or vertically).  
Mark each visited cell to avoid counting it again.

**12. How would you create a database schema for a university: colleges, students, professors? What tables, keys, and relationships?**

✅ **Answer:**

You can model this using a **relational schema** with the following tables:

■ **Tables and Keys:**

**1. College**

- college\_id (PK)
- college\_name
- location

**2. Department**

- department\_id (PK)
- department\_name
- college\_id (FK to College)

**3. Professor**

- professor\_id (PK)
- name
- email
- department\_id (FK to Department)

**4. Student**

- student\_id (PK)
- name
- email
- enrollment\_year
- college\_id (FK to College)

**5. Course**

- course\_id (PK)
- course\_name
- department\_id (FK to Department)

## 6. Enrollment

- student\_id (FK to Student)
- course\_id (FK to Course)
- semester
- grade
- *(Composite PK: student\_id + course\_id + semester)*

## 7. Teaching

- professor\_id (FK to Professor)
- course\_id (FK to Course)
- semester

📌 This ensures **1:N** (College → Students), **M:N** (Student ↔ Courses), and **M:N** (Professor ↔ Courses) relationships.

---

**13. Design an optimal schema to store event logs (like clicks/swipes) for high-velocity web traffic.**

✅ **Answer:**

For a high-volume, write-heavy log system:

🗄️ **EventLog Table (Wide, Denormalized, Partitioned)**

- event\_id (UUID or BIGINT, PK)
- user\_id (FK)
- event\_type (e.g., click, swipe)
- event\_time (timestamp)
- device\_type (mobile, desktop)
- location (optional - city/country)
- session\_id

- `page_url`
- `metadata` (JSON or MAP<string, string> — for extensibility)

#### 💡 Schema Design Tips:

- Store in **Delta Lake / Parquet** for compression and performance.
- Partition by **event\_date** or **event\_type** for efficient querying.
- Add **Z-ordering** on frequently filtered columns (e.g., `user_id` or `session_id`).

📌 Avoid joins during ingestion; keep schema flat.

---

**14. Given messy sales data, walk us through cleaning, transformation, and how you'd design reporting tables for business analytics.**

✅ Answer:

#### 🔪 Cleaning & Transformation Steps:

##### 1. Remove Duplicates:

- Use `dropDuplicates()` in PySpark or `ROW_NUMBER()` in SQL to keep the latest record.

##### 2. Handle Nulls:

- Fill missing region or salesperson using business logic or default values.
- Drop rows with missing `order_id` or `amount`.

##### 3. Data Type Fixing:

- Ensure dates are parsed correctly.
- Convert amount to Decimal or Double.

##### 4. Standardize Values:

- Convert categories to lowercase/title case.
- Format phone numbers or addresses.

#### 📁 Reporting Table Design:

##### 1. Fact\_Sales:

- `order_id` (PK)
- `customer_id`

- product\_id
- sales\_date
- amount
- discount
- region\_id
- salesperson\_id

## 2. Dim\_Date, Dim\_Customer, Dim\_Product, Dim\_Region, Dim\_Salesperson

📌 Use **star schema** for BI tools like Power BI or Tableau.

---

**15. Describe how you'd optimize a slow-performing JOIN query joining large tables. What indexing or partitioning strategies would you use?**

✅ **Answer:**

🔧 **Optimization Strategies:**

### 1. Broadcast Join:

- If one table is small (<10MB), use broadcast() to avoid shuffle.

### 2. Partition Pruning:

- Ensure tables are partitioned on join/filter columns.
- Example: sales\_data partitioned by region or sales\_date.

### 3. Clustered Index / Z-ordering:

- Use Z-ordering (Databricks) or clustered index on frequent filter/join columns.

### 4. Avoid Data Skew:

- If one key dominates (e.g., region\_id=1 in 90% rows), use **salting**.

### 5. Repartitioning:

- Repartition data before joining to ensure even distribution across tasks.

### 6. Filter Early:

- Apply WHERE conditions before joining to reduce row count.

### 7. Statistics & Caching:

- Update table stats so the optimizer chooses the right plan.

- Cache common lookups if reused.

📌 Always start with `.explain()` to check if joins are shuffled or optimized.